

Regression Analysis using GraphLab Create

Krishna Sridhar (GraphLab Team)

Overview

- Intro to Regression Analysis?
- Predictive Analytics
- Linear & Logistic regression!
- Interpreting Results
- Feature Engineering
- Feature Selection
- Bag of Trix
- Demo

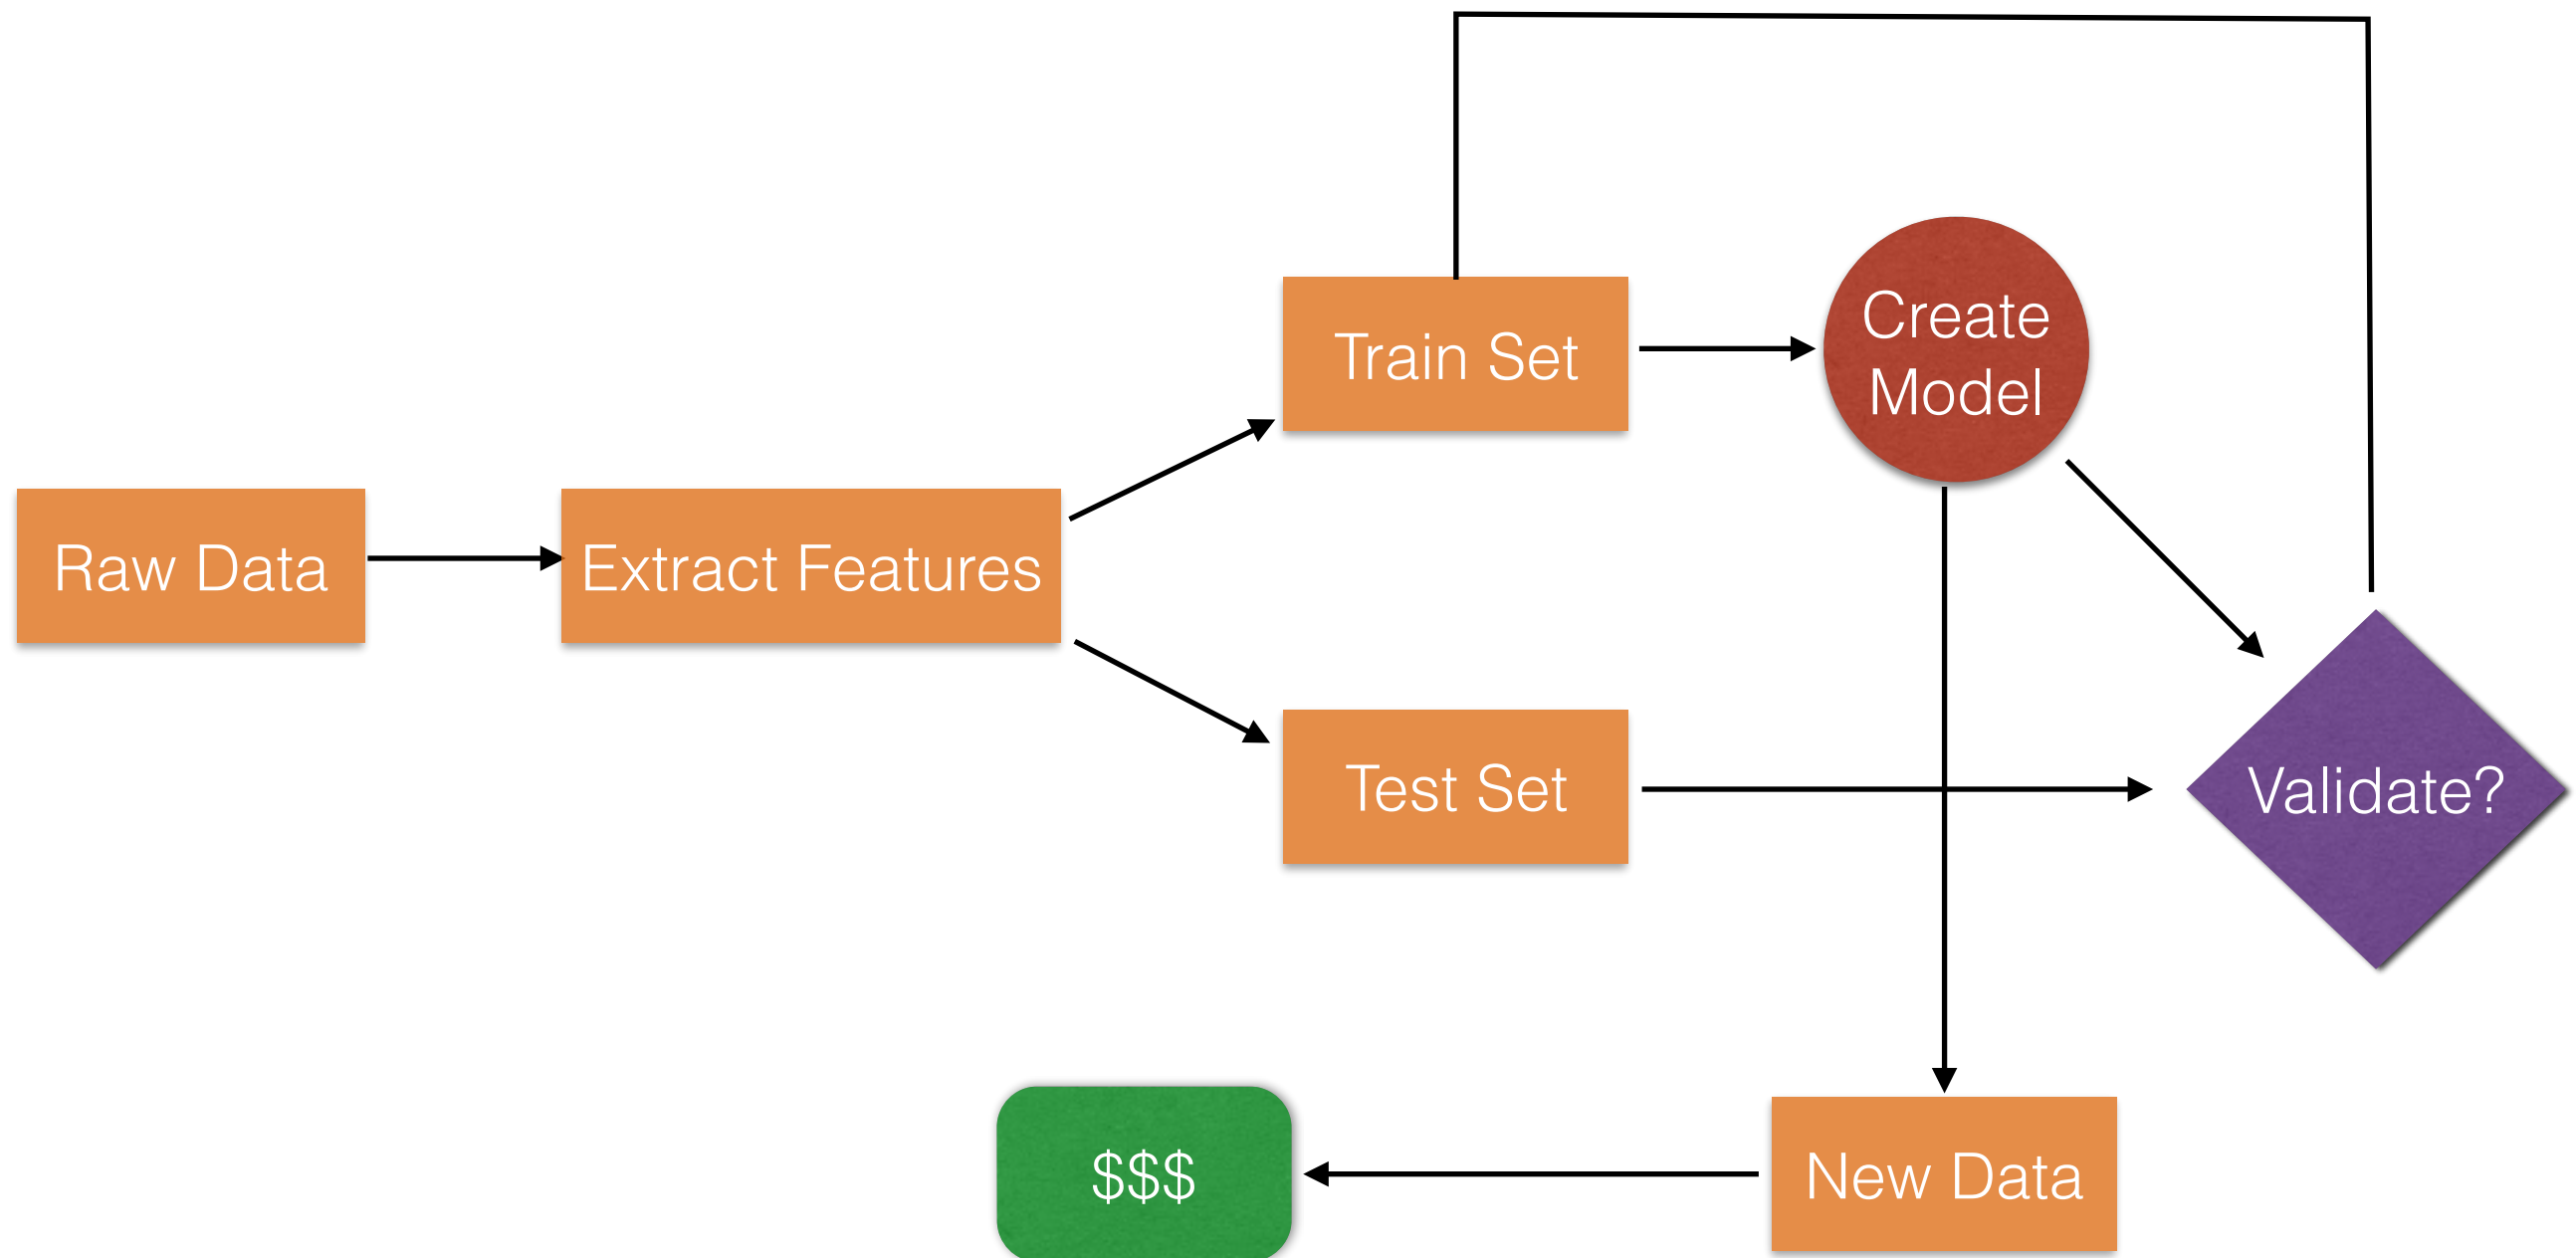
Predictive Analytics

The goal, in **predictive analytics** is to forecast or predict a target using historical data.

Examples

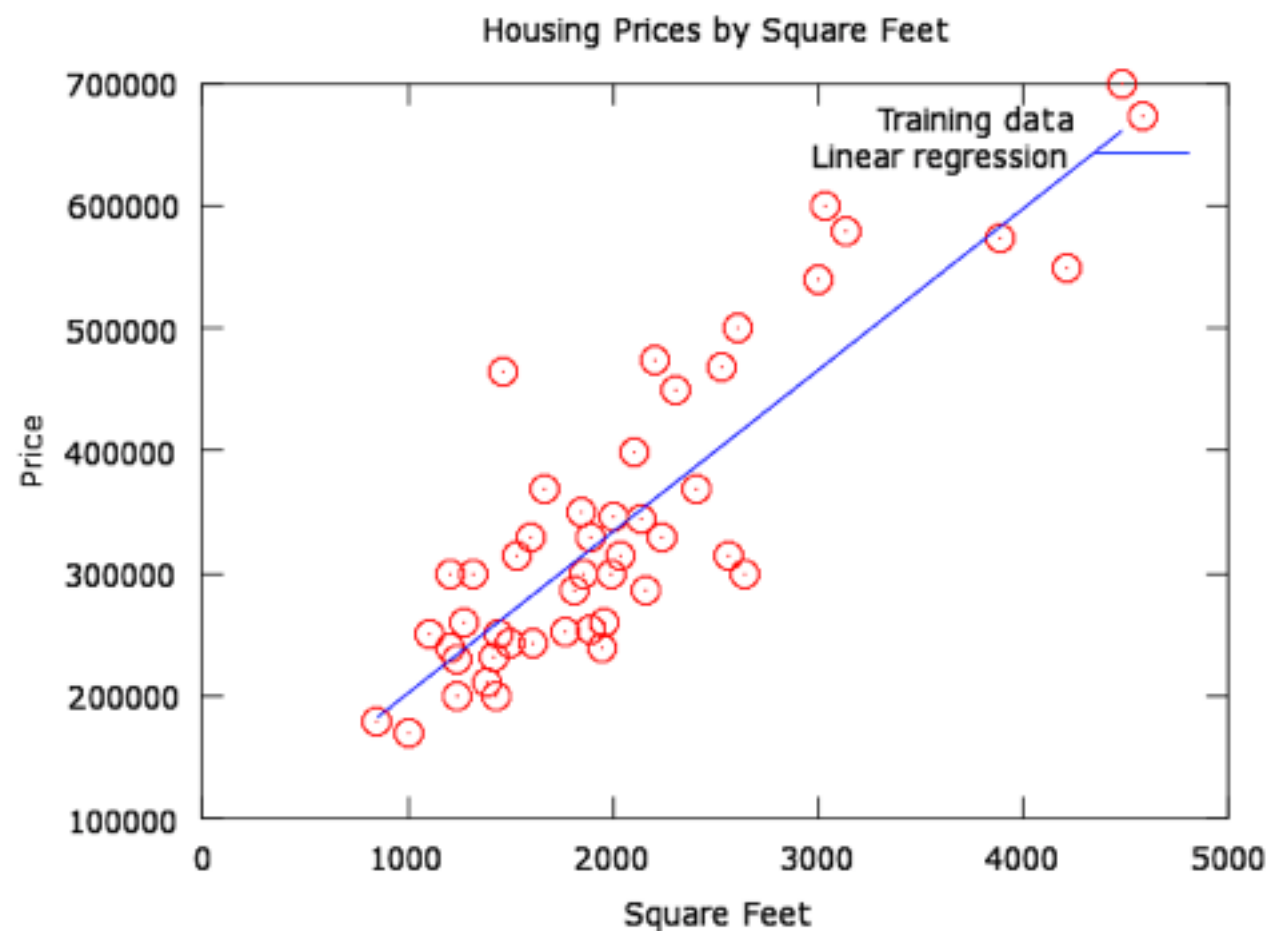
- What is the price of a house?
- Is a restaurant is good or not?
- Given a sentence, an I predict whether or not to say "That's what she said"
- Are the Heat going to win the NBA title?

Workflow




Predictive Analytics

Can I predict the price of a house given it's size?




Predictive Analytics

Can I predict if the Heat are going to win the NBA title?



Linear vs Logistic



Linear Regression

The diagram shows the linear regression equation $y_i = \beta_1 x_{i,1} + \beta_2 x_{i,2} + \dots + \beta_K x_{K,1} + \epsilon_i \quad \forall i \in \{1 \dots n\}$. The components are highlighted with colored circles: y_i is in a red circle, $x_{i,1}$ is in a purple circle, β_K is in a green circle, and ϵ_i is in a blue circle. Curved lines connect these circles to labels below: 'Target' for y_i , 'Features' for $x_{i,1}$, 'Coefficients' for β_K , and 'Noise/Error' for ϵ_i .

$$y_i = \beta_1 x_{i,1} + \beta_2 x_{i,2} + \dots + \beta_K x_{K,1} + \epsilon_i \quad \forall i \in \{1 \dots n\}$$

Target Features Coefficients Noise/Error

- **Target:** Binary (classification) or continuous (regression)
- **Features:** The set of attributes that effect how a decision/prediction is made.
- **Coefficients:** A measure of how the feature is correlated with the prediction.
- **Noise:** Nobody is prefect.

Training Models

The diagram shows the linear regression equation $y_i = \beta_1 x_{i,1} + \beta_2 x_{i,2} + \dots + \beta_K x_{K,1} + \epsilon_i$ for $\forall i \in \{1 \dots n\}$. The components are highlighted with colored circles and labeled with curved lines below them: y_i is in a red circle and labeled 'Target'; $x_{i,1}$ is in a purple circle and labeled 'Features'; β_K is in a green circle and labeled 'Coefficients'; and ϵ_i is in a blue circle and labeled 'Noise/Error'.

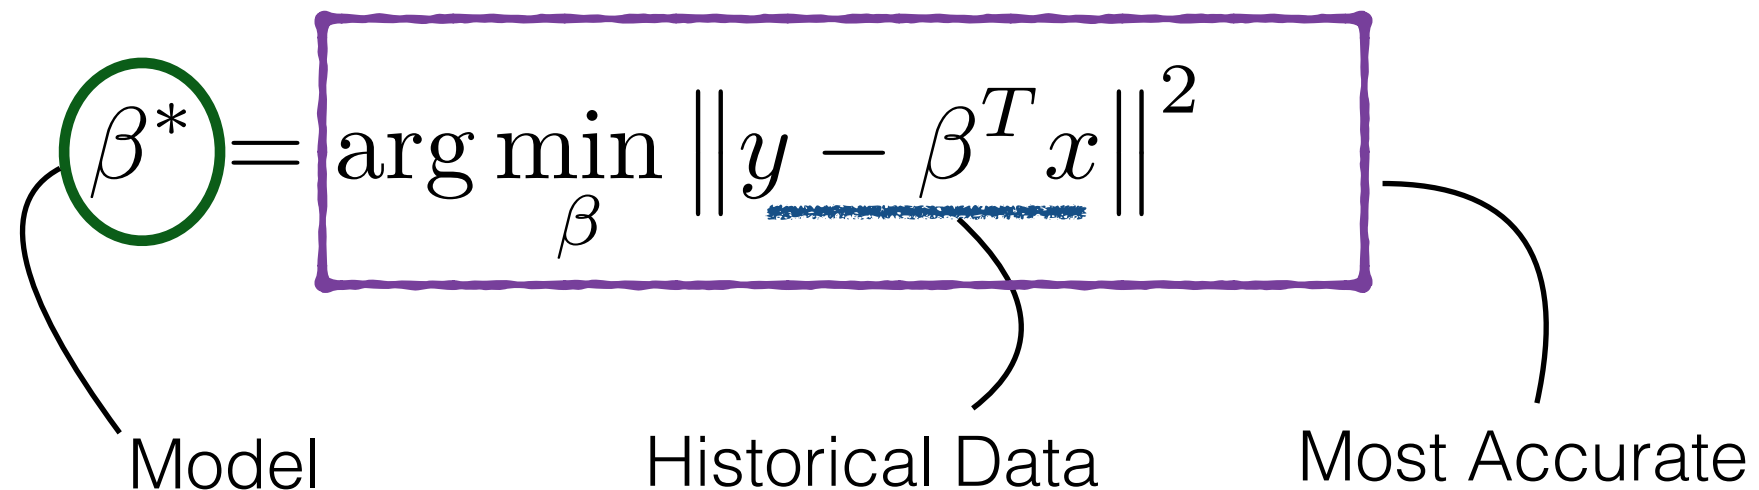
$$y_i = \beta_1 x_{i,1} + \beta_2 x_{i,2} + \dots + \beta_K x_{K,1} + \epsilon_i \quad \forall i \in \{1 \dots n\}$$

Target Features Coefficients Noise/Error

What is training?

Given some **historical data**, the process of model training tries to **find a model** which makes the **most accurate** predictions.

Training Models



The diagram shows the equation $\beta^* = \arg \min_{\beta} \|y - \beta^T x\|^2$. The term β^* is circled in green, with a line pointing to the label "Model". The entire right-hand side of the equation is enclosed in a purple rectangular box, with a line pointing to the label "Most Accurate". Inside this box, the term $\beta^T x$ is underlined in blue, with a line pointing to the label "Historical Data".

$$\beta^* = \arg \min_{\beta} \|y - \beta^T x\|^2$$

Model Historical Data Most Accurate

What is training?

Given some **historical data**, the process of model training tries to **find a model** which makes the **most accurate** predictions.

Training Models


The diagram shows the equation $\beta^* = \arg \min_{\beta} \|y - \beta^T x\|^2$. The term β^* is circled in green and labeled "Model". The entire equation is enclosed in a purple box, which is labeled "Historical Data". A bracket under the term $\beta^T x$ is labeled "Most Accurate".

$$\beta^* = \arg \min_{\beta} \|y - \beta^T x\|^2$$

Model Historical Data Most Accurate

How does training happen?

- Second order methods (Newton)
- First order methods (Gradient Descent, L-BFGS)
- Stochastic Methods (SGD)



Interpreting Results

$$\beta^* = \arg \min_{\beta} \|y - \beta^T x\|^2$$

Model Historical Data Most Accurate



Regularization

The diagram shows the equation $\beta^* = \arg \min_{\beta} \|y - \beta^T x\|^2 + \lambda \|x\|^2$. The term β^* is circled in green, and a line connects it to the label "Model". The entire equation is enclosed in a purple rectangular box, with a line connecting the box to the label "Most Accurate". The term $\beta^T x$ is underlined in blue, and a line connects it to the label "Historical Data".

$$\beta^* = \arg \min_{\beta} \|y - \beta^T x\|^2 + \lambda \|x\|^2$$


Model

Historical Data


Most Accurate

Regularization

A L1 regularization



B L2 regularization




L1 Penalty:


- Penalizes complex models.
- Creates models with lots of non-zero coefficients.
- Used for feature “selection”
- Suitable when you have more features than examples.
(Overdetermined problems)

Regularization

A L1 regularization



B L2 regularization



L2 Penalty:

- Avoid overfitting model for training data.
- Stabilize the estimates especially when there's collinearity in the data.
- Shrinks coefficients towards zero.

Demo Time!

Feature Engineering

GraphLab Create supports

- Numeric Features
- Categorical features
- Dictionaries (Sparse Features)
- Lists (Dense features)

Feature Engineering

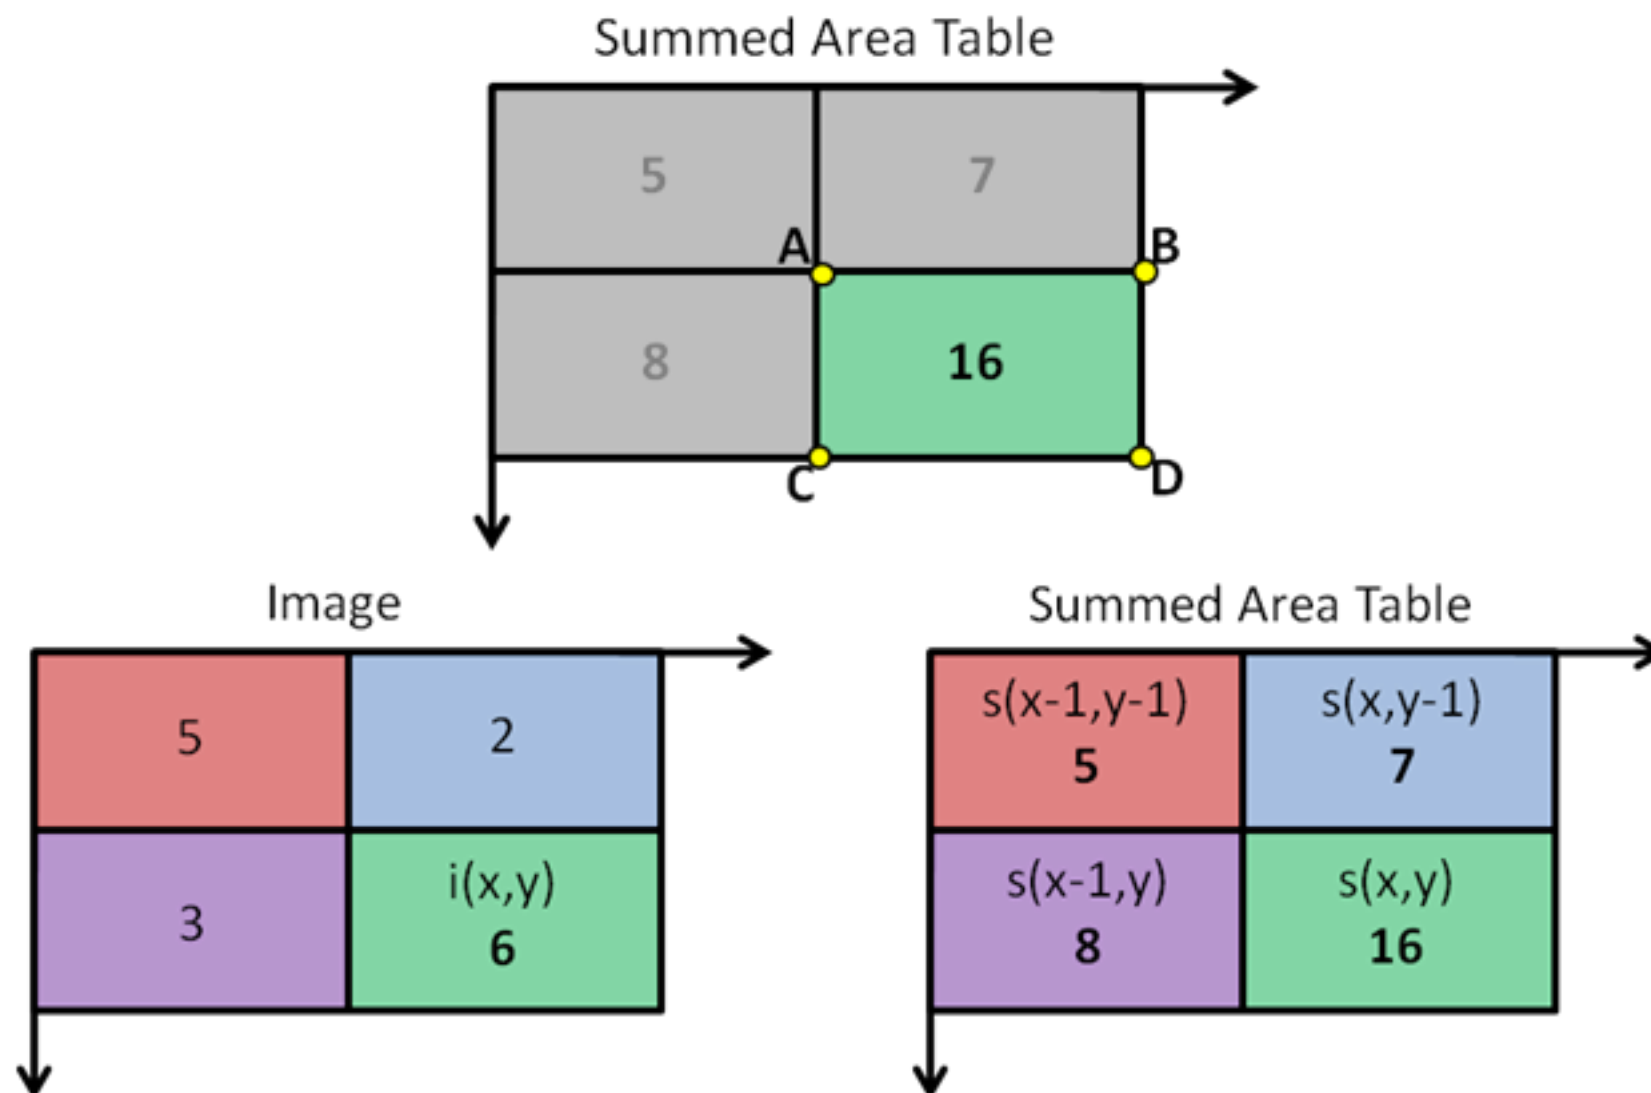
Categorical Features

- Eg: California (CA), Washington (WA), and Wisconsin (WI)

Category	Encoding
California	0 0
Washington	0 1
Wisconsin	1 0



Dense Vectors

Example: Viola-Jones Face detector



Dense Vectors

Example: Viola-Jones Face detector



Encode input features as a sequence of dense vectors

Sparse Features

Example: Bag of words

Doc. ID	Raw Text		Doc ID.	Bag of Words
1	In a bag of words, words corresponding ...		1	{'words': 2, 'in': 1 ...
2	I am a disco dancer ...		2	{'I': 1, 'am': 1 ...
3	That's what she said ...		3	{'that': 1, 'what': 1 ...

Encode input features as a sequence of sparse vectors

Feature Engineering

- Try interaction terms.
 - “Male Dog” is better than “Male” or “Dog”
- Plot residuals vs Individual features.
- Transformation on features (log, binning etc.)
- Nonlinear feature generation
 - Random features for large-scale kernel machines - Rahimi & Recht
 - Count-min sketch and feature hashing (see http://hunch.net/~jl/projects/hash_reps/)



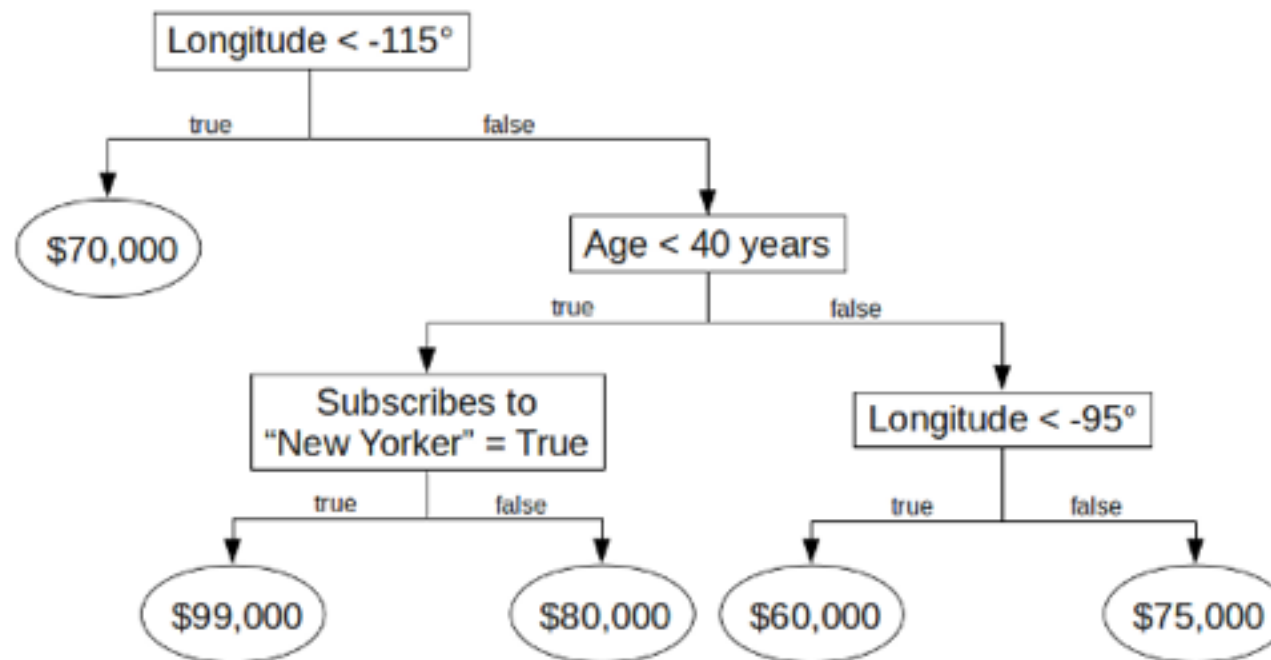
Vowpal Wabbit

Online Learning

- Learn as you go!
- Built to work with LOTS of features!
- **GraphLab Create's VW Wrapper Supports**
 - Linear, Logistic or Hinge Loss (SVM)
 - Matrix factorization and other modules using “command line args”.



Gradient Boosted Trees



GraphLab Create's Wraps "EXtreme Gradient Boosting"

- Tree learning for Regression and Classification
- Linear model / LASSO

Demo Time!